

Agenda

01. What are Diffusion Models?


02. Diffusion Model Architectures

03. Other Stuff

04. Research

1. What are Diffusion Models?


Use variational lower bound




- Diffusion models transform noise into sample from target data distribution.
- Model outputs are of the same shape (B, H, W, C) as inputs.

The origins of Diffusion Models comes from Langevin Dynamics.

Stochastic Gradient Descent $x_{j+1} = x_j - \alpha \nabla_x L(x), x_0 \sim \mathcal{N}(x)$




We can consider the density function as an optimization function



The goal is to find x that maximizes $p_{\text{data}}(x)$, exactly the same as in SGD. But for high dimensional data, like images.

$$x_{j+1} = x_j + \alpha \nabla_x \log p(x), \quad x_0 \sim \mathcal{N}(x|0, 1).$$

$\nabla_x \log p(x)$ – score function




But how to get a score function? We only have a dataset.

This is **the most important question** in diffusion models.

$$\|s_\theta(x) - \nabla_x \log p_{\text{data}}(x)\| - ?$$


We will not dive deeper into this question

$$\begin{aligned} & \|s_\theta(x) - \nabla_x \log p_{\text{data}}(x)\| \\ & \sim \text{tr}(\nabla_x s_\theta(x)) + \frac{1}{2} \|s_\theta(x)\|^2 \end{aligned}$$



$$\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\} \stackrel{\text{i.i.d.}}{\sim} p(\mathbf{x})$$

score
matching



$$\mathbf{s}_\theta(\mathbf{x}) \approx \nabla_{\mathbf{x}} \log p(\mathbf{x})$$

Langevin
dynamics




When sampling with Langevin dynamics, our initial sample is likely in low density regions


$$\mathbb{E}_{x \sim p_{\text{data}}} \|s_\theta(x) - \nabla_x \log p_{\text{data}}(x)\| \rightarrow \min_{\theta}$$

$$x_{j+1} = x_j + \alpha \nabla_x \log p(x), \quad x_0 \sim \mathcal{N}(x|0, 1).$$


Data density



Data scores




Estimated scores




Solution: Lets noise our data and learn the score for noised samples

$$\nabla_x \log p(x) \rightarrow \nabla_x \log p(x, \sigma),$$


Perturbed density




Perturbed scores



Estimated scores



We noise our data for different σ .



$$x_{j+1} = x_j + \alpha s_\theta(x_j, \sigma_j),$$
$$x_0 \sim \mathcal{N}(x|0, 1), \sigma_0 = 1, \sigma_N = 0.001$$

2. Diffusion Models Architectures




UNet?

Transformer?

SSM?


UNet

- The most popular architecture is old fellow **UNet**, familiar from the segmentation task.
- UNet is a stack of Residual Convolutional Blocks operating on different resolutions.
- Most diffusion models also include Attention blocks.
- Deeper layers have more channels and operate on smaller features maps.





What about transformer?


- A Diffusion model architecture could be built entirely from Transformer blocks.
- Vision Transformer (ViT) divides the input image into small patches, applies convolution, and processes the sequence like Language Models.
- Token positions are encoded using 2D positional embeddings.



Current State




Diffusion Transformer SD3/FLUX/...



2.1 Text encoding

- B : batch size
- $L_C = 77$: CLIP token length
- L_T : T5 token length
- d_0, d_1, d_T : embedding dims of CLIP-G, CLIP-L, and T5
- $d_T = 4096$

1. Text embedding



1.1 CLIP encoding

$$E_G \in \mathbb{R}^{B \times L_C \times d_0}, \quad E_L \in \mathbb{R}^{B \times L_C \times d_1},$$

1.2 Concat CLIP tokens

$$E_{\text{CLIP}} = [E_G; E_L] \in \mathbb{R}^{B \times L_C \times (d_0 + d_1)}$$


1.3 T5 Embeddings

$$E_{T5} \in \mathbb{R}^{B \times L_T \times d_T}$$

2.1 Text encoding

- B : batch size
- $L_C = 77$: CLIP token length
- L_T : T5 token length
- d_0, d_1, d_T : embedding dims of CLIP-G, CLIP-L, and T5
- $d_T = 4096$

1. Text embedding



1.4 Pad CLIP Channels to 4096

$$E'_{\text{CLIP}} = \text{Pad}(E_{\text{CLIP}}, (0, d_T - (d_0 + d_1))) \in \mathbb{R}^{B \times L_C \times d_T}$$

1.5 T5 Embeddings

$$E_{\text{final}} = [E'_{\text{CLIP}}; E_{T5}] \in \mathbb{R}^{B \times (L_C + L_T) \times d_T}$$

2.1 Text encoding


- B : batch size
- $L_C = 77$: CLIP token length
- L_T : T5 token length
- d_0, d_1, d_T : embedding dims of CLIP-G, CLIP-L, and T5
- $d_T = 4096$

2. Pooled text embedding

2.1 CLIP encodings

$$p_G \in \mathbb{R}^{B \times d_0}$$

$$p_L \in \mathbb{R}^{B \times d_1}$$




2.2 Concatenate Pooled CLIP Globals

$$p_{\text{CLIP}} = [p_G; p_L] \in \mathbb{R}^{B \times (d_0 + d_1)}$$

2.1 Text encoding

- B : batch size
- $L_C = 77$: CLIP token length
- L_T : T5 token length
- d_0, d_1, d_T : embedding dims of CLIP-G, CLIP-L, and T5
- $d_T = 4096$



$$c = E_{\text{final}} \in \mathbb{R}^{B \times (L_C + L_T) \times 4096}$$

$$y = p_{\text{CLIP}} \in \mathbb{R}^{B \times (d_0 + d_1)}$$

2.2 Noised Latent

$$z_t \in \mathbb{R}^{B \times C \times H \times W}$$


1. Patching

$$z_t^{(p)} \in \mathbb{R}^{B \times N \times (C \cdot P^2)}$$

$$N = \frac{H \cdot W}{P^2}$$

2. Linear Projection

$$h = z_t^{(p)} W_\ell + b_\ell, \quad W_\ell \in \mathbb{R}^{(C \cdot P^2) \times d_x}, \quad h \in \mathbb{R}^{B \times N \times d_x}$$




2.2 Noised Latent

3. Positional Embeddings

$$x = h + e_{\text{pos}} \in \mathbb{R}^{B \times N \times d_x}$$

$$e_{\text{pos}} \in \mathbb{R}^{N \times d_x}$$

$$x = z_t^{(p)} W_\ell + b_\ell + e_{\text{pos}}$$



2.3 Timestep

t : the diffusion timestep (scalar, e.g. $t \in [0, 1000]$)


1. Sinusoidal Encoding

$$\text{Sinusoidal}(t) = [\sin(t \cdot \omega_0), \cos(t \cdot \omega_0), \sin(t \cdot \omega_1), \cos(t \cdot \omega_1), \dots]$$

$$\omega_k = \frac{1}{10000^{2k/d_t}}, \quad k = 0, 1, \dots, \frac{d_t}{2} - 1$$

2. MLP

$$e_t = W_2 \text{SiLU}(W_1 \text{Sinusoidal}(t) + b_1) + b_2$$



$$e_t \in \mathbb{R}^{d_t}$$

2.3 Final pooled embedding


$$c = E_{\text{final}} \in \mathbb{R}^{B \times (L_C + L_T) \times 4096}$$

$$y = p_{\text{CLIP}} \in \mathbb{R}^{B \times (d_0 + d_1)}$$

$$e_t = W_2 \text{SiLU}(W_1 \sin(t) + b_1) + b_2 \quad \in \mathbb{R}^{B \times d_y}$$

$$h_y = W_y p_{\text{CLIP}} + b_y \quad \in \mathbb{R}^{B \times d_y}$$

$$y = e_t + h_y \quad \in \mathbb{R}^{B \times d_y}$$




2 Summary


$$(x, c, y)$$


$$x = \text{Patchify}(z_t)W_\ell + e_{\text{pos}} \in \mathbb{R}^{B \times N \times d_x}$$

$$c = [E_G; E_L; E_{T5}] \in \mathbb{R}^{B \times (L_C + L_T) \times 4096}$$

$$y = h_y + e_t$$








- **Normalization**
- **Modulation layers**
- **Linear layers**
- **Attention**

2.4 Modulation Layers

$$y \in \mathbb{R}^{B \times d_y}, \quad x \in \mathbb{R}^{B \times N \times d_x}.$$




$$[\alpha_x, \beta_x, \gamma_x, \delta_x, \epsilon_x, \zeta_x] = W_y \text{SiLU}(y) + b_y, \quad \alpha_x, \beta_x, \gamma_x, \delta_x, \epsilon_x, \zeta_x \in \mathbb{R}^{B \times d_x}.$$

2.4 Modulation Layers

$$[\alpha_x, \beta_x, \gamma_x, \delta_x, \epsilon_x, \zeta_x] = W_y \text{SiLU}(y) + b_y,$$

$$\tilde{x} = \alpha_x \odot \text{LayerNorm}(x) + \beta_x$$



2.4 Attention Layer

1. Linear

$$Q_x = x_{\text{mod}} W_Q^{(x)}, \quad K_x = x_{\text{mod}} W_K^{(x)}, \quad V_x = x_{\text{mod}} W_V^{(x)}$$


$$Q_c = c_{\text{mod}} W_Q^{(c)}, \quad K_c = c_{\text{mod}} W_K^{(c)}, \quad V_c = c_{\text{mod}} W_V^{(c)}$$

2. RMS Norm

$$Q = g_Q \odot \frac{Q'}{\sqrt{\frac{1}{d} \sum_k (Q'_k)^2 + \epsilon}}, \quad K = g_K \odot \frac{K'}{\sqrt{\frac{1}{d} \sum_k (K'_k)^2 + \epsilon}},$$

3. Concatenation

$$Q = [Q_x; Q_c], \quad K = [K_x; K_c], \quad V = [V_x; V_c].$$



4. Self-attention

$$A = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V,$$

2.5 Again modulation

1. Attention residual scaling


$$x_{\text{att}} = x + \gamma_x \odot A_x$$

2. LayerNorm + modulation

$$x_{\text{mlp-in}} = \delta_x \odot \text{LayerNorm}(x_{\text{att}}) + \epsilon_x$$

3. MLP

$$m_x = W_2^{(x)} \text{SiLU}(W_1^{(x)} x_{\text{mlp-in}} + b_1^{(x)}) + b_2^{(x)}$$



4. Gated residual update

$$x' = x_{\text{att}} + \zeta_x \odot m_x$$

$$[\alpha_x, \beta_x, \gamma_x, \delta_x, \epsilon_x, \zeta_x] = W_x(\text{SiLU}(W_y y))$$

$$[\alpha_c, \beta_c, \gamma_c, \delta_c, \epsilon_c, \zeta_c] = W_c(\text{SiLU}(W_y y))$$

$$x_m = \alpha_x \odot \text{LayerNorm}(x) + \beta_x,$$


$$c_m = \alpha_c \odot \text{LayerNorm}(c) + \beta_c$$

$$Q = [x_m W_Q; c_m W_Q], \quad K = [x_m W_K; c_m W_K],$$

$$V = [x_m W_V; c_m W_V] \quad A = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V$$

$$x' = x + \gamma_x \odot \text{Linear}(A_x) + \zeta_x \odot \text{MLP}(\delta_x \odot \text{LayerNorm}(A_x) + \epsilon_x)$$

$$c' = c + \gamma_c \odot \text{Linear}(A_c) + \zeta_c \odot \text{MLP}(\delta_c \odot \text{LayerNorm}(A_c) + \epsilon_c)$$



3. Image editing



(a) Input image



(b) “remove the thing from her face”




(c) “she is now taking a selfie in the streets of Freiburg, it’s a lovely day out.”



(d) “it’s now snowing, everything is covered in snow.”

3. Image editing



3. Personalization



Input images



in the Acropolis



swimming



sleeping




in a bucket



getting a haircut

3. Personalization



Input images (~3-5)

Reconstruction Loss



"A [V] dog"

Text → Image



3. Video generation


- Text-2-Video generation is quite similar to Text-2-Image.
- However, there are important differences:
 - Input data has now additional **temporal dimension**.
 - **Positional embeddings** have to account for this.
 - In addition to relevance and image quality, **motion consistency and smooth transition** between frames are crucial.
- Some approaches treat video clips as a single sequence, whereas make several predictions based on previous frames.



3. Video generation




4. Example of possible research




$$y = e_t + \cancel{h_y} \in \mathbb{R}^{B \times d_y}$$



w/o CLIP, long




w/o CLIP, short



with CLIP, short

Deviation from initial image (DreamSim) relative to number of tokens in prompt



$$\mathbf{y}(\mathbf{p}, t) \rightarrow \hat{\mathbf{y}}(\mathbf{p}, \mathbf{p}_+, \mathbf{p}_-, t) = \mathbf{y}(\mathbf{p}, t) + w \cdot (\mathbf{y}(\mathbf{p}_+, t) - \mathbf{y}(\mathbf{p}_-, t)).$$

- Short hair



Original generation



+ Long hair



\mathbf{p}^+ = Long hair | Modern car

\mathbf{p}^- = Short hair | Old car



- Old car



Original generation



+ Modern car 41

Ours, quality

FLUX, schnell



FLUX



Hi-Dream



SD3.5 Large



COSMOS



Original



a smiling banana wearing a bandana